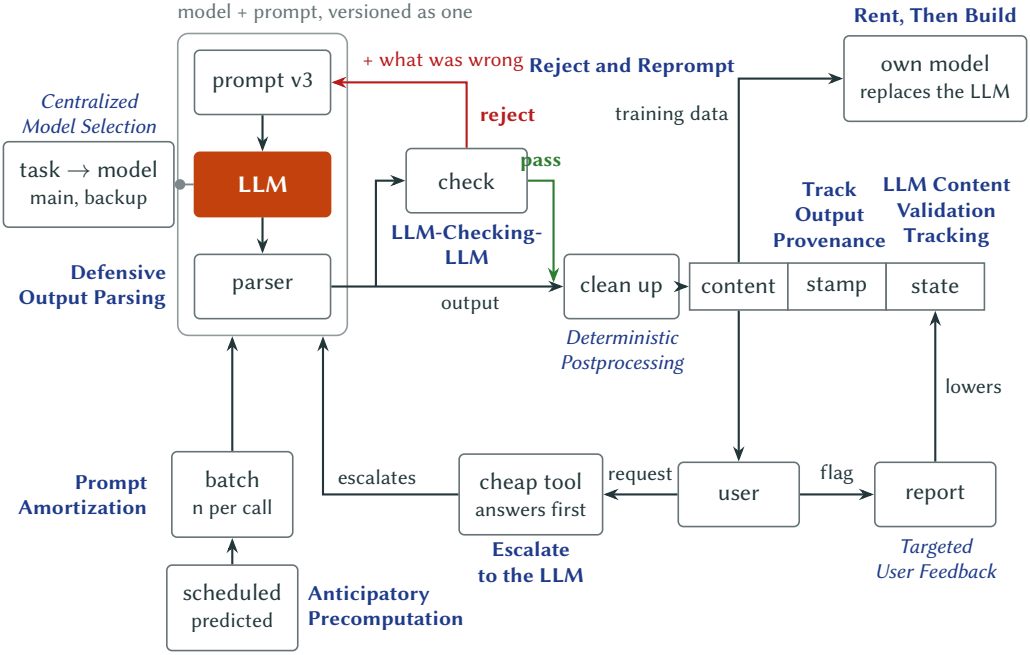
model + prompt, versioned as one
prompt v3
+ what was wrong
Reject and Reprompt
reject
Centralized
Model Selection
task → model
main, backup
LLM
check
pass
LLM-Checking-
LLM
Defensive
Output Parsing
parser
output
clean up
Deterministic
Postprocessing
content
stamp
state
Rent, Then Build
own model
replaces the LLM
training data
Track
Output
Provenance
LLM Content
Validation
Tracking
lowers
Prompt
Amortization
batch
n per call
escalates
cheap tool
answers first
request
user
flag
report
Escalate
to the LLM
Targeted
User Feedback
scheduled
predicted
Anticipatory
Precomputation